NEEDLE: A Live Benchmark of Web Search APIs for AI Agents

Ilya Gusev **Matthias Petri** **Andrey Styskin**
Keenable
{ilya.gusev,matthias,andrey}@keenable.ai

Abstract

NEEDLE is a live benchmark of web search APIs for AI agents. An agent sends such an API a text query and sees only the titles, URLs and snippets it returns. Agent benchmarks score the final answer, so they rarely test the search API alone. NEEDLE scores the API response directly. Hourly and daily runs build new queries from public sources for news, finance, research papers and law, and add a fixed stream of rare-word queries. We compare each engine with the best results that any tested engine found. Over eight weeks we test 26 search engines from 20 vendors. For 178 of 210 pairs of the 21 main engines, one engine scores reliably higher than the other on the same 12.4k queries. Searches by paper title hide large gaps: 23 engines find a paper from its title at recall@10 .71–.98, but from keywords in its body at only .08–.72. These results provide evidence that API comparisons should also test queries built from a paper's body.

1 Introduction

An **AI agent** is a language model that works in steps: it calls tools, reads their output and decides what to do next. Many agent tasks need facts from the web, such as a recent event, a company's quarterly income or a paper that another source cites. For these facts, the agent calls a **web search API**: a service that takes a text query and returns a short ranked list of results. Each result has a title, a URL and a snippet of page text. The agent sees only this list. It answers from the snippets or picks which pages to open. If the right page is not in the list, the agent cannot use it.

Many vendors sell such APIs, but a developer who must choose one has little public evidence. Agent benchmarks score the final answer, which mixes the search API with the model and the prompt (Section 2). Testing the API alone is hard.

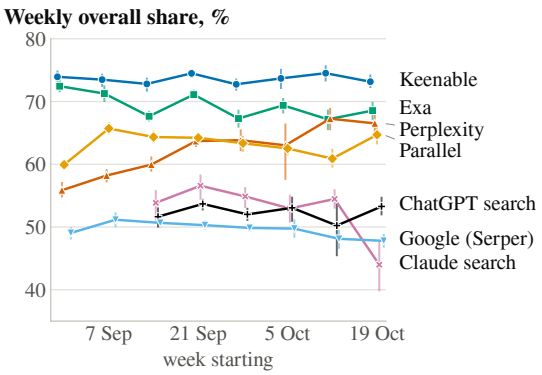


Figure 1: The top rank holds all eight weeks, but lower ranks change. Overall share by ISO week, 31 August–25 October 2026. ChatGPT and Claude search start on 14 September, and the pool grows from 18 to 25 engines. Points are query-weighted means, and bars are 95% run-level bootstrap intervals. Scores follow the dashboard rules.

The live web has no fixed document list or full relevance labels, and fixed test sets go stale and can leak into model training data (Sainz et al., 2023; Balloccu et al., 2024).

NEEDLE tests the search step alone. Each run sends the same new queries to every **engine**, which is one search API with fixed settings, and scores the results that each engine returns. The queries come from five tasks: news events, company finance facts, research papers, US law and rare-word queries from agent logs. For example, the scholar task looks for a paper with the query "gap rent regularity", whose words do not occur in the paper's title or abstract (Table 1). When a query has a known answer, such as a paper's arXiv id, we check whether the results contain it. Otherwise an LLM judge rates each result. We compare each engine with the **ceiling**: the best results that any tested engine found for that query.

We ran NEEDLE for eight weeks on 26 engines from 20 vendors. A **query-run** is one query issued

Task	Query	Reference answer or profile
news	Paramount Warner Bros takeover closes 2026	closing date, ephemeral event
finance, fact	nvidia number of employees	42,000, ±15%
finance, money	"servicenow" q3 fiscal 2024 operating income	$418M, ±0.1%
scholar, title	optimally selling heterogeneous objects by quantity	arXiv 2610.36244
scholar, body	gap rent regularity	same paper
legal, CFR	"modify or waive any" county committees authority	7 CFR 1417.2
rare	"Krasnogvardeysky District" Russia	judged, rare word has 7 pieces

Table 1: Sample generated queries. News and rare have no reference answer. For them, the judge profile specifies what a result must contain.

in one run. Our claims:

- **Most engine pairs have a clear winner.** For each pair of engines, we test whether one scores higher than the other on the same queries by more than chance explains. For 178 of the 210 pairs among the 21 main engines, it does, on 12.4k shared query-runs. A developer can rely on these orderings. For the other 32 pairs, the data cannot tell which engine is better (Section 4, Table 2).
- **Title searches hide large gaps.** When the query is a paper's title, 23 of 26 engines put the paper in the top ten for 71–98% of queries. When the query is keywords for a detail in the paper's body, the same engines find it for only 8–72%. An agent that follows up a detail it read elsewhere has no title to search for, so title tests can overstate how well an engine serves it (Section 4, Figure 2).
- **Rankings are stable from week to week, but not fixed.** Adjacent weekly rankings agree at Kendall's τ_b .81–.94, and the same engine leads all eight weeks. Lower ranks change, so a ranking from one week can misplace engines in the middle (Section 4, Figure 1).

2 Related Work

Live and API benchmarks. LiveBench and LiveSearchBench draw new items from dated sources but score model answers (White et al., 2025; Zhou et al., 2026a). NEEDLE scores the search results behind such answers. BrowseComp scores agents end to end (Wei et al., 2025). FreshStack and BrowseComp-Plus fix the corpus to isolate retrieval (Thakur et al., 2025; Chen et al., 2026), while NEEDLE keeps the live web. API comparisons use static questions under one agent (Selvam and Ghosh, 2026; Artificial Analysis, 2026), and web engine audits use fixed query sets (Lewandowski, 2015; Kirsten et al., 2026). NEEDLE draws new queries each run and scores the API without an agent.

Scorers and tasks. LLM assessors and simulated known-item queries can reproduce system rankings (Thomas et al., 2024; Upadhyay et al., 2025; Azzopardi et al., 2007; He et al., 2025). NEEDLE uses both and checks its judge against raters from other vendors. Domain benchmarks cover literature, caselaw and finance search on fixed sets (Ajith et al., 2024; Hou et al., 2025; Hu et al., 2026), while NEEDLE runs these tasks on new queries each day. Full text helps retrieval beyond abstracts (Lin, 2009), and our body queries test that gap on live APIs.

3 The Benchmark

We test 26 engines from 20 vendors, including five mode variants. Table 3 lists the tasks, and Appendix A the API settings and data sources.

3.1 Queries

News. Each hour an LLM writes 20 event queries from news feeds and Google Trends, rejecting evergreen items.

Finance. Queries ask for company facts from Wikidata (Vrandečić and Krötzsch, 2014) and GLEIF, quarterly 10-Q figures from SEC XBRL, or filings by description.

Scholar. Each recent arXiv or Europe PMC (Europe PMC Consortium, 2015) paper yields four paired known-item queries (Azzopardi et al., 2007): its normalized **title**, **body keywords** for a detail absent from title and abstract, a **body question** about that detail, and a **tip-of-tongue** description (Arguello et al., 2021) without names or exact values. All four must pass the filters.

Legal. CourtListener opinions yield caption queries, and eCFR sections yield queries based on verbatim spans. Filters reject citation leaks.

Rare. We sample a fixed public stream of English agent queries containing rare words (Ning et al., 2026; Li et al., 2026;

Engine	Overall share, %		Scholar	News	Finance	Rare	Legal	Failed, %	Latency
	mean	95% CI	R@10	nDCG@5	R@5	nDCG@5	R@5	requests	p_{50}, ms
ceiling (pooled)	*100*		*.901*	*1*	*.993*	*1*	*.971*		
Keenable	73.2	72.3–74.1	**.731**	.609	.903	**.508**	**.797**	0.0	135
Exa	68.5	67.4–69.4	.563	.599	**.922**	.476	.771	0.3	2,132
Perplexity	66.9	65.3–68.5	.578	.640	.895	.468	.674	0.0	433
Parallel	64.9	63.6–66.1	.587	.535	.907	.387	.734	0.8	2,660
Brave (LLM ctx)	64.6	63.4–65.6	.533	**.642**	.858	.454	.658	0.0	1,112
Linkup	57.6	56.1–58.9	.470	.567	.797	.400	.570	0.0	1,212
Brave	57.3	56.2–58.4	.488	.557	.769	.354	.621	0.0	487
Context	57.1	55.4–58.6	.598	.469	.782	.339	.576	0.0	866
DeepSeek	56.9	55.8–57.9	.271	.635	.818	.387	.676	0.0	551
You.com	55.8	54.7–56.9	.471	.553	.747	.343	.602	0.0	486
ChatGPT search†§	51.5	49.4–53.1	.300	.340	.809	.366	.697	11.1	43,831
Kagi§	51.3	49.9–52.6	.560	.301	.728	.275	.617	7.6	1,275
Firecrawl	50.7	49.2–52.0	.488	.286	.820	.265	.598	0.0	688
Octen	50.2	48.8–51.7	.307	.511	.730	.322	.586	0.0	197
Google (Serper)	47.7	46.8–48.6	.475	.262	.738	.255	.582	0.0	1,785
Claude search†§	45.1	42.7–47.6	.498	.276	.685	.195	.525	35.9	33,906
TinyFish§	44.5	43.0–45.9	.429	.256	.730	.220	.521	29.0	1,904
Bing (SearchAPI)	44.1	41.1–46.8	.416	.288	.675	.173	.586	0.4	2,008
Jina	43.8	42.4–45.1	.492	.210	.726	.163	.523	4.9	17,272
Tavily	42.6	41.5–43.7	.303	.303	.828	.284	.361	0.1	1,999
Ceramic	14.7	13.9–15.7	.001	.040	.441	.109	.137	0.1	99

Table 2: Results on 276 shared runs, 17–28 October 2026 (12.4k query-runs, 21 engines). The ceiling also includes variants. Task cells are absolute metrics, with the best in bold. Italics mark the ceiling. Failed counts all request failures. Latency is the per-run median on news. †Server-side search tool, so latency includes answer generation. §Account-side exclusions leave rank uncertain (Appendix C).

Zhou et al., 2026b). Table 1 gives query examples, and Appendix B gives the filters.

At least half of legal queries and finance money and filing queries carry quotes, `site:`, `after:` or `before:`. The client maps operators to native filters where available.

Task	Cadence	Queries/run	k	Metric
news	hourly	20	5	nDCG@5
finance	daily	120	5	recall@5
scholar	daily	up to 420	10	recall@10
legal	daily	40 + 0–20	5	recall@5
rare	daily	100	5	nDCG@5

Table 3: Benchmark tasks, run frequency and primary metrics. k is the number of search results requested per query. Query counts can fall below the targets after source and query filtering. Legal combines 40 CFR queries with caption queries for recent court opinions.

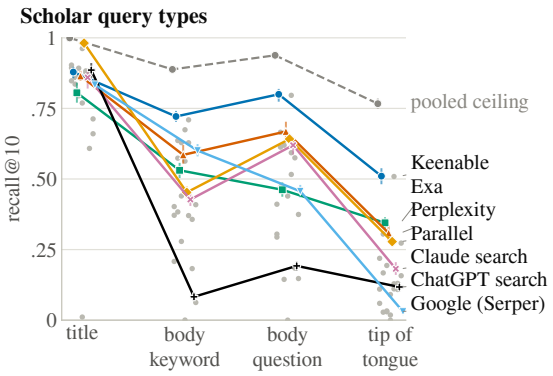


Figure 2: Engines that find papers from titles differ widely on body queries. Scholar recall@10 by query type: 31 runs, 29 September–28 October 2026, with 2,600–3,005 query-runs per line and type. Lines show seven main engines, grey dots the other 19 engines, and dashes the pooled ceiling. Points are query-weighted means, and bars are 95% run-level bootstrap intervals.

3.2 Scores

News and rare. GPT-6 Luna at high reasoning effort rates each distinct result from 0 to 4 using a shared query profile, with checks for relevance, coverage, freshness and literal operator compliance. The scale adapts the Needs Met levels of Google (2025). Ratings 0–4 map to gains $g = \{0, 0, .117, .667, 1\}$, putting high weight on strong results (Kanoulas and Aslam, 2009). Snippets have a 2,000-character cap.

For each query, the ceiling keeps the best rating per normalized URL across all engines, variants included, and sorts these ratings (Spärck Jones and van Rijsbergen, 1975; Zobel, 1998). With r_i the rating at rank i, we compute nDCG@5 (Järvelin

and Kekäläinen, 2002) as

$$\mathrm{DCG}_e(q) = \sum_{i=1}^{5} \frac{g(r_i)}{\log_2(i+1)},$$
$$\mathrm{nDCG}_e(q) = \frac{\mathrm{DCG}_e(q)}{\mathrm{IDCG}(q)}.$$

IDCG is the DCG of the top five pooled ratings. Missing ranks contribute zero. Repeated normalized URLs receive zero gain except on `site:` queries. We drop a query-run from all engines if any judgment fails or no result reaches rating 2. The task score is the mean over retained query-runs, and the judged ceiling has mean 1.

Scholar, legal and finance. Recall@k counts queries with a matching result in the top k. Scholar matches arXiv ids, DOIs and PMIDs. Legal matches citations, dockets and identifiers in URLs and snippets. Caselaw recall measures case identity, not necessarily the opinion page. Finance checks whether the snippets contain the answer, with a 0.1% money band and a 15% headcount band. The ceiling is a hit by any engine. Appendix F gives the matching rules and checks matcher error.

3.3 Overall share and failures

For task metric M_t, an engine's overall share is

$$S(e) = \frac{1}{5} \sum_{t=1}^{5} \frac{M_t(e)}{M_t(\mathrm{ceiling})}.$$

The tasks have equal weight. The ceiling is the best the tested engines found, not the best the web contains.

Table 2 uses dashboard scores from the 276 runs in which all 21 main engines took part. On finance, scholar and legal, we exclude account-side failures (our quota, rate limit, login or credit). Other failures score zero. On news and rare, every failure in this window scores zero. These exclusions remove up to 28% of an engine's query-runs.

Comparing two engines. Some queries are hard for every engine, so intervals for single engines can overlap even when one engine beats another on the same queries. We therefore bootstrap the difference $S(a) - S(b)$. Within each task, we resample whole runs 2,000 times, using the same runs for every engine (Efron and Tibshirani, 1993), and recompute the overall shares. Each difference gets a normal-approximation interval with Bonferroni correction for all $\binom{21}{2} = 210$ pairs at a 5% familywise level (Carterette, 2012). An interval that excludes zero orders that pair. The intervals measure sampling uncertainty, not scorer error (Appendix C).

3.4 Judge validation

We sample 100 news and 100 rare queries, with six distinct top-five URLs per query (1,200 results). Two Claude Opus 5.5 raters, Gemini 4 Argon and Grok 4.7 grade these with engine names hidden and a common written guideline. Claude raters also see the first 3,000 page characters. The Claude raters agree at quadratic-weighted $\kappa = .94$ (Cohen, 1968). Gemini and Grok agree with their consensus at .92 and .90. Agreement with the production judge is .77–.79 (Table 7). Gemini and Grok also grade 100 complete query pools (5,698 results from all 21 engines) from snippets only. Each gives engine-rank Spearman .94 against the production judge, and none separates Keenable from Exa. These checks measure model agreement, not human validation (Appendix E).

4 Results

Most engine pairs have a clear winner, but not all. For 178 of 210 engine pairs, the adjusted interval for the score difference excludes zero, so one engine scores reliably higher than the other. We cannot order the remaining 32 pairs from this sample. An unresolved pair does not mean equal scores. The overall lead also depends on the task mix: scholar contributes 3.7 points of Keenable's 4.7-point lead over Exa. Without scholar, the lead is 1.2 points (Appendix C).

Title searches hide large gaps. Over 31 scholar runs, 23 of 26 engines reach recall@10 .71–.98 from titles. The same 23 span .08–.72 on body keywords (Figure 2 and Table 6). Keenable leads at .72, followed by its realtime mode (.71) and Kagi (.67). Appendix D compares engines on common papers.

Rankings are stable from week to week, but not fixed. Adjacent weekly rankings agree at Kendall's τ_b (Kendall, 1945) .81–.94. Keenable leads all eight weeks, but lower ranks change (Figure 1). Appendix C checks changes in the engine pool, account failures and dispatch times.

Limitations

Conflict of interest. The authors' organization builds an evaluated engine. The code, prompts,

clients and per-query outputs are public so readers can audit the claims. All engines receive the same queries, snippet cap, request pacing and judge. Each keeps one setting throughout the main window. We cannot tell whether a vendor changed its index for the benchmark or for other reasons.

Evidence and judge. Scores measure the API response, not the fetched page. Short snippets can hide relevant content (Turpin et al., 2009; Kazai et al., 2022), and identifier matching can miss a target without its id. On 400 sampled results, page text changes 25% of grades. News and rare depend on one judge. All audit raters are LLMs, and none provides human validation, which LLM labels cannot replace (Faggioli et al., 2023; Clarke and Dietz, 2025). Shared rules can cause shared errors. Query-term stuffing can fool judges (Alaofi et al., 2024), and vendor and model biases can affect scores (Balog et al., 2025; Soboroff, 2025). The judge shares a vendor with ChatGPT search, and some audit raters share one with Claude search. LLM evaluators can favor their own outputs (Panickssery et al., 2024).

Matcher error. Finance table figures can match the wrong period when snippets omit column headers. In a fresh LLM audit, 5 of 335 table-reading hits contradict the answer. Another 95 show the right figure without a period label, which the snippet alone cannot confirm. The pooled audit cannot show each engine's error rate. Caselaw recall credits pages about the case, not just the opinion. Appendix F gives the finance checks.

Coverage and weights. Queries are English and US-centric. Equal task weights do not estimate agent traffic, and the lead depends strongly on scholar. Default API tiers need not be each vendor's best. ChatGPT and Claude search return sources through answer-generating tools, so their quality scores are comparable but their latency and cost are not. The ceiling depends on the engine pool (Buckley et al., 2007), so compare shares only within a common window.

Failures and timing. Excluding account-side failures assumes those requests would score like the rest. Without that assumption, several middle ranks remain uncertain. News and rare also count our credit lapses as zeros in the main window. Engines process queries at different speeds, so slow engines see news later. A five-minute dispatch filter leaves the news order unchanged. Vendor redirect links can break URL matching.

Repetition and reproducibility. Registry facts and the fixed rare stream can repeat. Public sources and queries permit tuning or special handling of benchmark traffic. The studied runs archive queries, reference answers, results and judgments, but lack API publication dates and each run's code revision. Stored ratings reproduce the tables, but re-judging cannot exactly reproduce the inputs. Appendix A gives the source terms.

References

Anirudh Ajith, Mengzhou Xia, Alexis Chevalier, Tanya Goyal, Danqi Chen, and Tianyu Gao. 2024. LitSearch: A retrieval benchmark for scientific literature search. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 15068–15083.

Marwah Alaofi, Paul Thomas, Falk Scholer, and Mark Sanderson. 2024. LLMs can be fooled into labelling a document as relevant: best café near me; this paper is perfectly relevant. In *Proceedings of the 2024 Annual International ACM SIGIR Conference on Research and Development in Information Retrieval in the Asia Pacific Region (SIGIR-AP)*, pages 32–41.

Jaime Arguello, Adam Ferguson, Emery Fine, Bhaskar Mitra, Hamed Zamani, and Fernando Diaz. 2021. Tip of the tongue known-item retrieval: A case study in movie identification. In *Proceedings of the 2021 Conference on Human Information Interaction and Retrieval (CHIIR)*, pages 5–14.

Artificial Analysis. 2026. Artificial analysis search index: Same agent, different search. Launched August 18, 2026; accessed November 3, 2026.

Leif Azzopardi, Maarten de Rijke, and Krisztian Balog. 2007. Building simulated queries for known-item topics: An analysis using six European languages. In *Proceedings of the 30th Annual International ACM SIGIR Conference on Research and Development in Information Retrieval*, pages 455–462.

Simone Balloccu, Patrícia Schmidtová, Mateusz Lango, and Ondřej Dušek. 2024. Leak, cheat, repeat: Data contamination and evaluation malpractices in closed-source LLMs. In *Proceedings of the 18th Conference of the European Chapter of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 67–93.

Krisztian Balog, Donald Metzler, and Zhen Qin. 2025. Rankers, judges, and assistants: Towards understanding the interplay of LLMs in information retrieval evaluation. In *Proceedings of the 48th International ACM SIGIR Conference on Research and Development in Information Retrieval*, pages 3865–3875.

Chris Buckley, Darrin Dimmick, Ian Soboroff, and Ellen M. Voorhees. 2007. Bias and the limits of pooling for large collections. *Information Retrieval*, 10(6):491–508.

Benjamin A. Carterette. 2012. Multiple testing in statistical analysis of systems-based information retrieval experiments. *ACM Transactions on Information Systems*, 30(1):4:1–4:34.

Zijian Chen, Xueguang Ma, Shengyao Zhuang, Ping Nie, Kai Zou, Sahel Sharifymoghaddam, et al. 2026. BrowseComp-Plus: A fair and disentangled evaluation benchmark for deep search agents. In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 22349–22370.

Charles L. A. Clarke and Laura Dietz. 2025. LLM-based relevance assessment still can't replace human relevance assessment. In *Proceedings of the International Workshop on Evaluating Information Access (EVIA 2025), a Satellite Workshop of the NTCIR-18 Conference*, pages 1–5.

João Coelho, Jingjie Ning, Jingyuan He, Kangrui Mao, Abhijay Sai Paladugu, Pranav Setlur, et al. 2026. DeepResearchGym: A free, transparent, and reproducible sandbox for deep research. In *Proceedings of the 2026 International ACM SIGIR Conference on Innovative Concepts and Theories in Information Retrieval (ICTIR)*, pages 34–43.

Jacob Cohen. 1968. Weighted kappa: Nominal scale agreement provision for scaled disagreement or partial credit. *Psychological Bulletin*, 70(4):213–220.

Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. 2019. BERT: Pre-training of deep bidirectional transformers for language understanding. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 4171–4186.

Bradley Efron and Robert J. Tibshirani. 1993. *An Introduction to the Bootstrap*, volume 57 of *Monographs on Statistics and Applied Probability*. Chapman & Hall, New York.

Europe PMC Consortium. 2015. Europe PMC: a full-text literature database for the life sciences and platform for innovation. *Nucleic Acids Research*, 43(D1):D1042–D1048.

Guglielmo Faggioli, Laura Dietz, Charles L. A. Clarke, Gianluca Demartini, Matthias Hagen, Claudia Hauff, et al. 2023. Perspectives on large language models for relevance judgment. In *Proceedings of the 2023 ACM SIGIR International Conference on the Theory of Information Retrieval (ICTIR)*, pages 39–50.

Google. 2025. *Search Quality Rater Guidelines: General Guidelines*. Version of September 11, 2025.

Yifan He, To Eun Kim, Fernando Diaz, Jaime Arguello, and Bhaskar Mitra. 2025. Tip of the tongue query elicitation for simulated evaluation. In *Proceedings of the 48th International ACM SIGIR Conference on Research and Development in Information Retrieval*, pages 3398–3407.

Abe Bohan Hou, Orion Weller, Guanghui Qin, Eugene Yang, Dawn Lawrie, Nils Holzenberger, et al. 2025. CLERC: A dataset for U.S. legal case retrieval and retrieval-augmented analysis generation. In *Findings of the Association for Computational Linguistics: NAACL 2025*, pages 7913–7928.

Liang Hu, Jianpeng Jiao, Jiashuo Liu, Dongyuan Mutu, Yanle Ren, Zhoufutu Wen, et al. 2026. FinSearchComp: Towards a realistic, expert-level evaluation of financial search and reasoning. In *The Fourteenth International Conference on Learning Representations (ICLR)*.

Kalervo Järvelin and Jaana Kekäläinen. 2002. Cumulated gain-based evaluation of IR techniques. *ACM Transactions on Information Systems*, 20(4):422–446.

Armand Joulin, Edouard Grave, Piotr Bojanowski, Matthijs Douze, Hervé Jégou, and Tomas Mikolov. 2016. FastText.zip: Compressing text classification models. *Preprint*, arXiv:1612.03651.

Armand Joulin, Edouard Grave, Piotr Bojanowski, and Tomas Mikolov. 2017. Bag of tricks for efficient text classification. In *Proceedings of the 15th Conference of the European Chapter of the Association for Computational Linguistics: Volume 2, Short Papers*, pages 427–431.

Evangelos Kanoulas and Javed A. Aslam. 2009. Empirical justification of the gain and discount function for nDCG. In *Proceedings of the 18th ACM Conference on Information and Knowledge Management (CIKM)*, pages 611–620.

Gabriella Kazai, Bhaskar Mitra, Anlei Dong, Nick Craswell, and Linjun Yang. 2022. Less is less: When are snippets insufficient for human vs machine relevance estimation? In *Advances in Information Retrieval: 44th European Conference on IR Research (ECIR 2022), Part II*, pages 153–162.

M. G. Kendall. 1945. The treatment of ties in ranking problems. *Biometrika*, 33(3):239–251.

Elisabeth Kirsten, Jost Grosse Perdekamp, Qinyuan Wu, Mihir Upadhyay, Krishna P. Gummadi, and Muhammad Bilal Zafar. 2026. Characterizing web search in the age of generative AI. In *Findings of the Association for Computational Linguistics: ACL 2026*, pages 10827–10848.

Dirk Lewandowski. 2015. Evaluating the retrieval effectiveness of web search engines using a representative query sample. *Journal of the Association for Information Science and Technology*, 66(9):1763–1775.

Zhuofeng Li, Dongfu Jiang, Xueguang Ma, Haoxiang Zhang, Ping Nie, Yuyu Zhang, Kai Zou, Jianwen Xie, Yu Zhang, and Wenhu Chen. 2026. OpenResearcher: A fully open pipeline for long-horizon deep research trajectory synthesis. In *Proceedings of the 2026 Conference on Empirical Methods in Natural Language Processing*, pages 908–937.

Jimmy Lin. 2009. Is searching full text more effective than searching abstracts? *BMC Bioinformatics*, 10:46.

Jingjie Ning, João Coelho, Yibo Kong, Yunfan Long, Bruno Martins, João Magalhães, Jamie Callan, and Chenyan Xiong. 2026. Agentic search in the wild: Intents and trajectory dynamics from 14m+ real search requests. In *Proceedings of the 49th International ACM SIGIR Conference on Research and Development in Information Retrieval*, pages 1404–1415.

Arjun Panickssery, Samuel R. Bowman, and Shi Feng. 2024. LLM evaluators recognize and favor their own generations. In *Advances in Neural Information Processing Systems 37 (NeurIPS 2024)*, pages 68772–68802.

Oscar Sainz, Jon Ander Campos, Iker García-Ferrero, Julen Etxaniz, Oier Lopez de Lacalle, and Eneko Agirre. 2023. NLP evaluation in trouble: On the need to measure LLM data contamination for each benchmark. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pages 10776–10787.

Sriram Selvam and Anneswa Ghosh. 2026. Similar accuracy, unequal evidence: Search APIs as decision surfaces for tool-using agents. In *Proceedings of the 1st Workshop on Grounding Language Models: Learning Faithfully and Efficiently (GroundLM 2026)*, pages 606–622.

Ian Soboroff. 2025. Don't use LLMs to make relevance judgments. *Information Retrieval Research*, 1(1):29–46.

Karen Spärck Jones and C. J. van Rijsbergen. 1975. Report on the need for and provision of an 'ideal' information retrieval test collection. British Library Research and Development Report 5266, Computer Laboratory, University of Cambridge.

Nandan Thakur, Jimmy Lin, Sam Havens, Michael Carbin, Omar Khattab, and Andrew Drozdov. 2025. FreshStack: Building realistic benchmarks for evaluating retrieval on technical documents. In *Advances in Neural Information Processing Systems (NeurIPS), Datasets and Benchmarks Track*.

Paul Thomas, Seth Spielman, Nick Craswell, and Bhaskar Mitra. 2024. Large language models can accurately predict searcher preferences. In *Proceedings of the 47th International ACM SIGIR Conference on Research and Development in Information Retrieval*, pages 1930–1940.

Andrew Turpin, Falk Scholer, Kalervo Järvelin, Mingfang Wu, and J. Shane Culpepper. 2009. Including summaries in system evaluation. In *Proceedings of the 32nd Annual International ACM SIGIR Conference on Research and Development in Information Retrieval*, pages 508–515.

Shivani Upadhyay, Ronak Pradeep, Nandan Thakur, Daniel Campos, Nick Craswell, Ian Soboroff, and Jimmy Lin. 2025. A large-scale study of relevance assessments with large language models using UMBRELA. In *Proceedings of the 2025 International ACM SIGIR Conference on Innovative Concepts and Theories in Information Retrieval (ICTIR)*, pages 358–368.

Denny Vrandečić and Markus Krötzsch. 2014. Wikidata: A free collaborative knowledgebase. *Communications of the ACM*, 57(10):78–85.

Jason Wei, Zhiqing Sun, Spencer Papay, Scott McKinney, Jeffrey Han, Isa Fulford, et al. 2025. BrowseComp: A simple yet challenging benchmark for browsing agents. *Preprint*, arXiv:2504.12516.

Colin White, Samuel Dooley, Manley Roberts, Arka Pal, Benjamin Feuer, Siddhartha Jain, et al. 2025. LiveBench: A challenging, contamination-limited LLM benchmark. In *International Conference on Learning Representations (ICLR)*.

Yonghui Wu, Mike Schuster, Zhifeng Chen, Quoc V. Le, Mohammad Norouzi, Wolfgang Macherey, et al. 2016. Google's neural machine translation system: Bridging the gap between human and machine translation. *Preprint*, arXiv:1609.08144.

Heng Zhou, Ao Yu, Yuchen Fan, Li Kang, Jianing Shi, Yongting Zhang, et al. 2026a. An automated pipeline for provably retrieval-dependent benchmark construction over dynamic knowledge. In *Findings of the Association for Computational Linguistics: EMNLP 2026*, pages 40633–40648.

Yuqi Zhou, Sunhao Dai, Changle Qu, Liang Pang, Jun Xu, and Ji-Rong Wen. 2026b. Learning to retrieve from agent trajectories. In *Proceedings of the 49th International ACM SIGIR Conference on Research and Development in Information Retrieval*, pages 2544–2555.

Justin Zobel. 1998. How reliable are the results of large-scale information retrieval experiments? In *Proceedings of the 21st Annual International ACM SIGIR Conference on Research and Development in Information Retrieval*, pages 307–314.

A API Settings and Data Sources

All engines receive the same query and requested result count. Each engine processes one request at a time, and engines run concurrently. Requests have a 30-second timeout and up to three attempts, with exponential back-off from 1 to 30 seconds on

HTTP 429, 5xx and timeouts. Short result lists score as returned. Snippets have a 2,000-character cap.

The generators and judge use GPT-6 Luna (`openai/gpt-6-luna`) through OpenRouter, with temperature 0 and high reasoning effort. Table 4 lists the search clients and settings. Appendix G gives all prompts.

The rare stream draws from DeepResearchGym (Coelho et al., 2026; Ning et al., 2026) (59%, CC BY 4.0), OpenResearcher (Li et al., 2026) (39%, MIT) and LRAT-Train (Zhou et al., 2026b) (1%, no stated license). Source terms govern query texts. Our code is MIT, and our CC BY 4.0 license covers results, judgments and metadata, not source text. LRAT contributes 17 of 1,200 main-window rare queries, and removing them changes no mean rare nDCG@5 by more than .003. Queries can contain personal information. Release contacts accept private removal requests. API output records remain subject to vendor terms.

B Query Filters

Scholar samples five domains (computer, physical, life, social and health sciences) and three age bands (7 days, 30 days, 1 year), targeting seven accepted papers per cell. Each paper must pass all four filters:

- **Title:** at least four words, plus a digit, acronym, hyphenated term or five content tokens.
- **Body keywords:** at least three words and two content tokens absent from title and abstract.
- **Body question:** at least eight words and the same novelty test as body keywords.
- **Tip-of-tongue:** at least twelve words and at most two content tokens shared with the title, with no quotes, decimal, five-digit number or coined title/abstract term.

An instrumented run accepts 323 of 572 candidate papers (56%). The effect of rejecting other papers on engine order is untested.

News draws from 124 feeds and Google Trends RSS for 52 US geographies. Prompts reject evergreen items and request two-to-six-word event queries. Legal captions come from 14 federal courts on CourtListener. CFR queries contain one verbatim two-to-eight-word span but no citation or target section number. Only the quoted variant retains quotation marks. Rare queries must pass fastText English language ID (Joulin et al., 2017, 2016) and contain a word split into at least five BERT WordPieces (Wu et al., 2016; Devlin et al., 2019).

C Statistical Checks and Ranking Sensitivity

Paired intervals. The main window has 12 scholar, 232 news, 10 finance, 12 rare and 10 legal runs. For pair (a, b), let $\widehat{\Delta}_{ab} = S(a) - S(b)$ and let s_{ab} be the standard deviation of its 2,000 bootstrap differences. The adjusted interval is

$$\widehat{\Delta}_{ab} \pm 3.683\, s_{ab}.$$

The interval is a normal approximation, not a percentile bootstrap interval. The implemented critical value is slightly more conservative than 3.675, the two-sided Bonferroni value for 210 comparisons at familywise level .05. Individual engine intervals in Table 2 instead use the 2.5th and 97.5th bootstrap percentiles. Their overlap is not a test of the difference between two engines.

Dependence and missing requests. We repeat the comparisons by resampling days, and by clustering exact query text (target paper for scholar) across runs. The latter checks repeated targets, which day resampling alone does not address. We also restrict the data to query-runs with no account-side failure in any main engine. Of the 178 pairs whose primary adjusted intervals exclude zero, 164 also exclude zero under both day resampling and this matched-support check.

Table 5 separates the failure-policy change from the change in query support. Per-query files cover 274 of the 276 runs, and matching retains 4,048 of 12,313 query-runs. With one account-failure policy, Claude search rises from 45.1 to 51.4 overall, and matching changes it to 52.6. The matched order has $\tau_b = .89$ against the main table, with Keenable and Exa still first and second. Its difference intervals use 1,000 paired run-bootstrap resamples. If excluded requests instead score zero or at the ceiling, ranks 1–10 stay fixed under both extremes, but Claude search and TinyFish can occupy ranks 11–20. Missing requests therefore limit some rank claims.

Normalization, weights and timing. Removing an engine from the pool changes its overall share by at most 0.8 points. Dropping any one task preserves the leader but changes the Keenable–Exa gap from 1.2 to 6.4 points. The lead's size is sensitive to task choice even when the leader stays the same.

Vendor	Endpoint	Variant or non-default parameters
Keenable	api.keenable.ai	default, realtime variant
Exa	api.exa.ai	`type=auto`, instant and lite variants
Google	google.serper.dev (Serper)	default
Bing	searchapi.io (SearchAPI)	engine `bing`
Brave	api.search.brave.com	web search, LLM-context endpoint
Parallel	api.parallel.ai	default processor, turbo variant
Perplexity	api.perplexity.ai	search API, fast variant
Tavily	api.tavily.com	default
Linkup	api.linkup.so	default
TinyFish	api.search.tinyfish.ai	default
Octen	api.octen.ai	default
Ceramic	api.ceramic.ai	default
You.com	ydc-index.io	default
Firecrawl	api.firecrawl.dev	search endpoint, default
Kagi	kagi.com/api/v1	default
Jina	s.jina.ai	default
Context	api.context.dev/v1	default
DeepSeek	api.deepseek.com	`/v1/tools/search`, default
Claude search	api.anthropic.com	`claude-sonnet-5`, `web_search`
ChatGPT search	api.openai.com	`gpt-5.5`, `web_search`, `search_context_size=low`

Table 4: Search clients and settings by vendor. Clients live in `src/needle/shared/search/` and also implement the native query filters.

Engine	Historical	Consistent policy	Matched	Δ matched	95% CI for Δ
Keenable	73.2	73.2	74.7	+1.6	+0.4, +2.7
Exa	68.5	68.6	67.5	−1.1	−2.3, +0.2
Perplexity	66.9	66.9	66.2	−0.7	−2.5, +1.1
Parallel	64.9	65.1	63.0	−2.1	−3.5, −0.7
Brave (LLM ctx)	64.6	64.6	67.1	+2.5	+1.1, +3.9
Linkup	57.6	57.6	59.3	+1.7	−0.1, +3.8
Brave	57.3	57.3	59.1	+1.8	+0.2, +3.2
Context	57.1	57.1	56.2	−0.9	−3.0, +1.1
DeepSeek	56.9	56.8	58.0	+1.1	−0.6, +2.9
You.com	55.8	55.8	57.4	+1.6	+0.3, +2.9
ChatGPT search	51.5	52.4	51.9	−0.5	−2.0, +0.8
Kagi	51.3	52.1	52.3	+0.1	−0.8, +1.2
Firecrawl	50.7	50.7	50.1	−0.5	−1.9, +0.9
Octen	50.2	50.2	49.1	−1.1	−2.7, +0.6
Google (Serper)	47.7	47.7	47.4	−0.3	−1.5, +0.9
Claude search	45.1	51.4	52.6	+1.2	−0.4, +2.9
TinyFish	44.5	45.1	44.7	−0.4	−2.0, +1.2
Bing (SearchAPI)	44.1	44.1	48.3	+4.1	+1.5, +7.5
Jina	43.8	43.8	43.2	−0.5	−2.2, +1.2
Tavily	42.6	42.6	42.5	−0.1	−1.2, +1.3
Ceramic	14.7	14.7	14.7	−0.1	−1.1, +1.2

Table 5: Overall shares in percent. Historical repeats Table 2. Consistent policy excludes account-side failures on every task and scores other failures zero, using 12,313 available query-runs. Matched restricts these to 4,048 query-runs with no account failure in any main engine. Δ is Matched minus Consistent policy, in percentage points, with a paired run-bootstrap interval.

Among the 15 engines present in all eight plotted weeks, adjacent rankings have $\tau_b = .85$–$.94$. Fixing the news pool to those 15 changes no sampled weekly order. Restricting news comparisons to requests sent within five minutes of the first engine retains 96% of scored pairs, leaves the order unchanged and moves no mean nDCG@5 by more than .010. These checks address changing engine participation and dispatch lag.

D Scholar Query Types and Paired Differences

Table 6 reports every engine and its sample size, rather than only the seven lines in Figure 2. In the main window, Keenable's recall lead over Exa is .38 for 7-day papers, .10 for 30-day papers and .03 for one-year papers. In health sciences, the lead is negative ($-.10$). The advantage is not uniform

Engine	Title	Body keyword	Body question	Tip of tongue	Runs	Query-runs
ceiling (pooled)	*1*	*.888*	*.938*	*.767*	*31*	*12,020*
Parallel	**.981**	.453	.642	.278	31	12,020
Kagi	.962	.674	.640	.091	30	11,599
Bing (SearchAPI)	.903	.402	.411	.060	31	12,020
Parallel Turbo	.894	.208	.294	.110	31	12,020
Brave (LLM ctx)	.893	.440	.620	.194	31	12,020
ChatGPT search	.886	.083	.192	.118	30	11,016
You.com	.883	.370	.593	.129	31	12,020
Brave	.881	.383	.613	.167	31	12,020
Keenable	.879	**.721**	**.800**	**.510**	31	12,020
Linkup	.878	.357	.575	.130	21	7,898
Keenable realtime	.877	.709	.796	.509	31	12,020
Perplexity	.865	.585	.668	.309	30	11,636
Jina	.862	.629	.470	.030	31	11,716
Perplexity Fast	.860	.574	.666	.305	30	11,636
Claude search	.860	.427	.620	.182	27	10,400
Context	.856	.628	.516	.111	31	12,020
Exa Lite	.848	.371	.305	.219	1	420
Firecrawl	.836	.636	.472	.028	31	12,020
Google (Serper)	.834	.601	.457	.030	31	12,020
Exa	.805	.531	.462	.345	31	12,020
Tavily	.795	.279	.144	.033	31	12,020
TinyFish	.769	.563	.440	.018	31	3,472
Exa Instant	.713	.409	.374	.273	31	12,020
Octen	.660	.063	.147	.157	31	12,020
DeepSeek	.608	.183	.184	.109	12	4,820
Ceramic	.011	.000	.000	.000	31	12,007

Table 6: Scholar recall@10 by query type, 31 runs, 29 September–28 October 2026. The ceiling has 3,005 query-runs per type, and the last column sums all four types. The 31 runs contain 332–420 queries each, totaling 12,020 rather than 31×420: 420 is a target, not a fixed sample size. Run counts show scored support, and Exa Lite has only one run. Counts also fall after account failures. Bold marks the best engine per type (`bucket_table.py`).

across age or domain.

For paired query-type comparisons, we retain 2,629 paper-runs from 29 runs where Keenable, its realtime mode, Kagi, Exa and ChatGPT search all have scores for all four types (10,516 query-runs). Account-side failures exclude the paper from all five engines, and other failures score zero. We use stored identifier hits and 2,000 paired run-bootstrap resamples.

Body-keyword recall is .154 lower than title recall for Keenable and .802 lower for ChatGPT search. On body keywords, Keenable leads its realtime mode by .012, Kagi by .048, Exa by .199 and ChatGPT search by .644. The 95% intervals of all these paired differences exclude zero. They use common papers, unlike the available-case means in Table 6, and do not measure identifier-matcher error. Full contrasts, intervals and replay inputs are in `audit/scholar_paired.json` and its companion JSONL file.

E Judge Sensitivity Checks

Audit protocol. The 200 queries in Section 3.4 come from 17–29 October runs with at least 19 engines. Each Claude rater uses a separate random order in 24 sessions of 50 items. Gemini and Grok grade one item per call at temperature 0. All use the same written guideline. Raters assign mean grades of 1.87–1.97 versus the production judge's 1.60, so agreement does not imply identical grade calibration.

Quadratic-weighted κ	News	Rare	All
Claude A vs. B	.94	.94	.94
Judge vs. Claude consensus	.72	.80	.77
Gemini vs. consensus	.91	.92	.92
Grok vs. consensus	.90	.88	.90
Gemini vs. judge	.74	.81	.78
Grok vs. judge	.75	.81	.79

Table 7: Agreement on 1,200 blinded results, 600 per task. Consensus is the rounded mean of the two Claude page grades.

Evidence length. The staged snippet-versus-page check changes 25% of 400 grades, but the mean grade change is only +.05, and its interval includes zero. Positive and negative changes can cancel in the mean.

On 803 news results from eight queries, a Claude

Opus 5.5 judge grades both full snippets and versions capped at 200 characters. The resulting engine orders agree at $\tau_b = .73$, and the mean absolute share change is .084. Exa and Keenable lose, while Bing (SearchAPI) and Google (Serper) gain. The sample does not distinguish the top two at either length. Snippet length is thus part of the measured API quality, not a controlled property of the underlying page.

Literal quotes. Of 1,200 rare query-runs, 504 contain quotes. On unquoted queries, Exa, Perplexity and Brave (LLM ctx) lead, and Keenable is fourth. Replacing the rare score with its unquoted-only value gives overall $\tau_b = .96$ and a 3.3-point Keenable lead over Exa. Literal phrase matching therefore affects the rare ranking even when the overall order changes little. These checks use LLM judgments, not human labels.

F Finance Matching and Checks

Matching rules. Money must fall within 0.1% of the reference value, and headcounts within 15%. A scale word fixes a money figure's scale. Bare figures, including parenthesized losses, can mean thousands or millions only when the preceding 200 characters contain a currency or financial cue. Years, percentages and multiples do not qualify. Diluted EPS also needs "EPS", "per share" or "earnings per". Money gets no LLM backstop. Other fields can receive an upgrade, never a revoked hit. The full rules are in `src/needle/finance/canon.py`.

Why the narrow money band? An LLM audit of 399 results checks company, period and line item. Money-hit precision without table reading is .93 at the 0.1% band and .88 at 2%: the wider band admits more wrong periods and related line items. Money backstop upgrades have precision .36, which motivates excluding them. These are model judgments, not human gold labels.

Unseen-query check. A separate sample has 307 unseen queries (146 money, 161 other), tested on 17 engines. An LLM checks the first matching result for 335 table-reading hits, 100 exact-band money hits and 50 other-field hits, using at least the full 2,000-character input. Exact-band precision is .93. Of 335 table-reading hits, 235 confirm the correct figure and period (.70), 95 omit period labels and five contradict the answer. Counting unresolved hits as wrong gives precision .70, and counting only contradictions gives .99. Confirmation rates differ by engine: for example, 13/43 for Firecrawl and 42/44 for Exa. Pooled precision cannot establish each engine's error rate.

Effect on reported ranks. Discounting each engine's main-window table-reading hits by its fresh unconfirmed fraction lowers finance recall by up to 6.0 points, and overall ranks 1–10 stay fixed. Dropping table reading entirely lowers finance recall by 0.1–7.5 points, with finance $\tau_b = .84$ and overall .97. Overall ranks 1–3 stay fixed. The check bounds one scoring choice, not all matcher error. The fresh sample has more money queries than production and lacks the non-money backstop, so it cannot replace the main table.

G Prompts

These Jinja templates are verbatim copies of the prompts in `src/needle/*/prompts/`. The code fills the braces at run time. Scholar title queries use no prompt.

Query generation. `projection.jinja`: News queries from feed items.

```
Today's date: {{ today }}

You are a search query generator for a search engine
      evaluation system. Generate a single search
     query that targets a specific news event
     happening today (within the past few hours)
     related to this recently-published article.

Requirements:
- Keyword-style query, not a natural language
     question
- 2-6 words, terse and direct
- The query MUST reference the specific event from
     the article, not just the person/team/topic in
     general (e.g. 'Lakers trade deadline' NOT just
     'Lakers')
- Include the year or date where it helps
     disambiguate
- Do NOT copy article titles or URLs verbatim
- Do NOT start the query with 'what' or 'why' -- use
      keyword style
- Always respond in English. If the article is not
     in English, respond with an English query
     anyways
- If the article's ideal answer would be the same a
     month from now (evergreen explainer, how-to,
     recipe collection, best-of list, gift guide,
     review of a stable product, historical
     retrospective, opinion piece), respond with
     exactly NO_NEWS_EVENT

Source kind: {{ source_kind }}
Title: {{ title }}
Summary: {{ summary }}
URL: {{ url }}

Respond with ONLY the query string (or NO_NEWS_EVENT
     ). No quotes, no explanation, no JSON.
```

`trends_projection.jinja`: News queries from Google Trends items.

```
Today's date: {{ today }}

You are a search query generator for a search engine
      evaluation system. Generate a single search
```

```
    query that targets a specific news event
    happening today (within the past few hours)
    related to this trending topic.

Requirements:
- Keyword-style query, not a natural language
    question
- 2-6 words, terse and direct
- The query MUST reference the specific event from
    the news articles, not just the person/team/
    topic in general (e.g. 'Lakers trade deadline'
    NOT just 'Lakers')
- Include the year or date where it helps
    disambiguate
- Do NOT copy article titles or URLs verbatim
- Do NOT start the query with 'what' or 'why' -- use
     keyword style
- Always respond in English. If the topic is not in
    English, respond with an English query anyways
- If the news articles do not describe a specific
    current event, respond with exactly
    NO_NEWS_EVENT

Trending topic: {{ topic }}

News context:
{{ news_block }}

Respond with ONLY the query string (or NO_NEWS_EVENT
    ). No quotes, no explanation, no JSON.
```

filing_query.jinja: Finance filing queries by description.

```
You generate a search query for a financial-search
    benchmark. The query must lead a web search
    engine to ONE specific SEC filing -- without
    ever citing its accession number.

You are given the filing company, form type, filing
    date, and an excerpt of the filing text.
    Produce ONE keyword-style search query anchored
     on a SPECIFIC, DISTINCTIVE detail from the
    excerpt: a named transaction or counterparty,
    an acquired business, a named facility or
    product, a precise unusual amount, an executive
     change, a named lawsuit or settlement.

Requirements:
- Keyword-style, 4-10 words total. Not a natural-
    language question. Do not start with what/why/
    how.
- Include the company name (or its common short form
    ) OUTSIDE the quotes.
- Wrap EXACTLY ONE distinctive span of 2-6
    consecutive words in double quotes, copied
    CHARACTER-FOR-CHARACTER from the excerpt. Pick
    wording specific to THIS filing -- not
    securities boilerplate (not "forward-looking
    statements", "pursuant to section", "securities
     and exchange commission", "quarterly report",
    "check mark").
- Do NOT include the filing's accession number, CIK,
     file number, or any URL.
- Always respond in English.
- If the excerpt is only boilerplate (cover page,
    checkboxes, table of contents) with no
    distinctive detail, respond with exactly
    NO_DISTINCT_QUERY.

Company: {{ company }}
Form: {{ form }}
Filed: {{ filed }}
Excerpt:
{{ text }}

Respond with ONLY the query string (or
    NO_DISTINCT_QUERY). No explanation, no JSON.
```

code_query.jinja: Legal CFR queries.

```
You generate a search query for a legal-search
    benchmark. The query must lead a web search
    engine to ONE specific section of the US Code
    of Federal Regulations -- without ever naming
    that section.

You are given the section's heading and its full
    text. Produce ONE keyword-style search query
    built from the section's SUBSTANCE: the
    specific obligation, prohibition, threshold,
    exemption, or defined term it establishes.

Requirements:
- Keyword-style, 4-10 words total. Not a natural-
    language question. Do not start with what/why/
    how.
- Wrap EXACTLY ONE distinctive span of 2-6
    consecutive words in double quotes, copied
    CHARACTER-FOR-CHARACTER from the section text (
    not the heading). Pick a span whose wording is
    regulatory and specific -- a defined term, an
    operative phrase, an unusual threshold -- not
    boilerplate that appears across the entire CFR
    (e.g. not "it shall be unlawful", "for purposes
     of this", "any person who").
- The rest of the query (outside the quotes) adds
    context words naming the subject matter, so the
     query works even on engines that ignore quotes.

- Do NOT include the section number, part number,
    title number, "CFR", "USC", "§", or the word "
    section" followed by a number. The task is to
    find the section from its content, not its
    citation.
- Do NOT copy the heading verbatim as the whole
    query.
- Always respond in English.
- If the section text is too generic or
    administrative to single it out (cross-
    references, addresses, OMB control numbers,
    severability boilerplate), respond with exactly
     NO_DISTINCT_QUERY.

Heading: {{ heading }}
Section text:
{{ text }}

Respond with ONLY the query string (or
    NO_DISTINCT_QUERY). No explanation, no JSON.
```

body_query.jinja: Scholar body-keyword queries.

```
You generate a search query for a scholarly-search
    benchmark. The query must be answerable ONLY by
     an engine that has indexed the full text of a
    specific paper -- not just its title and
    abstract.

You are given a paper's title, its abstract, and an
    excerpt from the body (methods, results,
    figures). Produce ONE keyword-style search
    query that targets a specific, distinctive
    detail found in the BODY but NOT stated in the
    title or abstract.

The query must be anchored on a GLOBALLY DISTINCTIVE
     term that THIS paper introduces, names, builds
    , or measures -- one that, on its own, already
    narrows the web to a handful of papers:
- a method, model, tool, software package, dataset,
    or benchmark the paper introduces or uses by
    name (e.g. "Toom's rule", "GC-MS carvacrol", "
    REDCap")
- a named study site, cohort, organism, cell line,
    reagent, or instrument
- a precise measured value that is itself unusual (e
    .g. "carvacrol 80.43%"), NOT a round number or
    a vague range

Ask yourself: would this query, typed into a
    scholarly search engine, plausibly return THIS
    paper near the top? If the query is built only
    from generic field vocabulary (e.g. "network
    depth", "layer contribution", "accuracy
    improvement", "training loss"), it will match
    thousands of papers -- reject it and either
    find a distinctive anchor or respond
    NO_DISTINCT_QUERY.

Shared configuration values are NOT distinctive:
    token counts, context-window sizes, batch sizes
    , parameter counts, learning rates, epochs, and
```

```
      similar hyperparameters recur across thousands
      of papers (e.g. "128 tokens", "128K context",
     "7B parameters", "batch size 32", "512 hidden
     dim"). Never anchor on these. Anchor on the
     specific NAME the paper gives its own artifact,
      or on a genuinely unusual measured result.

Requirements:
- Keyword-style, 3-8 words. Not a natural-language
     question. Do not start with what/why/how.
- At least one word must be a globally distinctive
     proper noun, named entity, or unusual precise
     value -- not generic jargon and not a shared
     hyperparameter/config value.
- Do NOT lead with, or build the query around, a
     percentage range, a round number, or a common
     configuration value.
- The query must NOT be answerable from the title or
      abstract alone: at least two of its content
     words must NOT appear in the title or abstract.
- Anchor on what THIS paper contributes -- NOT on
     prior work it cites and NOT on a well-known
     shared result. Specifically:
  - Do NOT use a citation or reference: no "et al",
     no author-surname strings (e.g. "Smith et al
     2024" or "Johnstone Güdel"). Cited names point
     at OTHER papers.
  - Do NOT anchor on an eponymous or textbook
     formula, theorem, law, transform, or
     distribution that predates and is shared across
      the field (e.g. "Cardano's formula", "Navier-
     Stokes", "Kermack-McKendrick", "Fourier
     transform").
- Do NOT reference the paper's own figures, tables,
     sections, theorems, or supplementary material
     by number (e.g. "Table 3", "Figure 2", "Theorem
      7.1", "Supplementary Table") -- those are
     internal pointers, not searchable content.
- Do NOT copy the title. Do NOT include the paper's
     arXiv id, DOI, or a URL.
- Always respond in English.
- If ONE distinctive multi-word span (3-6 words)
     appears VERBATIM in the body excerpt, wrap
     exactly that span in double quotes so exact-
     match engines can phrase-match it; copy it
     character-for-character from the body. Leave
     the rest of the query unquoted. If no
     distinctive verbatim span exists, use no quotes.

- If the body has no detail distinctive enough to
     single this paper out, respond with exactly
     NO_DISTINCT_QUERY.

Title: {{ title }}
Abstract: {{ abstract }}
Body excerpt:
{{ body }}

Respond with ONLY the query string (or
     NO_DISTINCT_QUERY). No explanation, no JSON.
```

`clue_query.jinja`: Scholar body-question queries.

```
You generate a search query for a scholarly-search
     benchmark. The query is a natural-language
     question a researcher might type to re-find a
     specific paper from precise factual clues.

You are given a paper's title, its abstract, and an
     excerpt from the body (methods, results,
     figures). Write ONE full-sentence question
     built from 2-3 precise, distinctive facts found
      in the BODY but NOT stated in the title or
     abstract -- the experimental setup, a named
     cohort, site, or instrument, or a specific
     measured result.

Example of the style (for a different paper): "In a
     German study of 145 cosmetic and personal care
     samples, what percentage contained NDMA and
     what was the maximum concentration found?"

Requirements:
- One question, 12-35 words, natural English prose.
     Not a keyword list.
- The clues must be precise enough that the question
      plausibly identifies THIS paper: include at
     least one named entity, specific setup detail,
     or unusual measured value from the body.
- At least two of the question's content words must
     NOT appear in the title or abstract.
- Anchor on what THIS paper did -- NOT on prior work
      it cites: no "et al", no author-surname
     strings.
- Do NOT reference the paper's own figures, tables,
     sections, theorems, or supplementary material
     by number.
- Do NOT copy the title. Do NOT include the paper's
     arXiv id, DOI, or a URL. No quotation marks.
- Always respond in English.
- If the body has no facts distinctive enough to
     single this paper out, respond with exactly
     NO_DISTINCT_QUERY.

Title: {{ title }}
Abstract: {{ abstract }}
Body excerpt:
{{ body }}

Respond with ONLY the question (or NO_DISTINCT_QUERY
     ), on a single line. No explanation, no JSON.
```

`tot_query.jinja`: Scholar tip-of-tongue queries.

```
You generate a search query for a scholarly-search
     benchmark. The query is a tip-of-the-tongue
     request: how a researcher describes a paper
     months after reading it, when the details have
     gone fuzzy.

You are given a paper's title, its abstract, and an
     excerpt from the body. Write ONE hedged, half-
     remembered description (1-2 sentences) of the
     paper's setup and main finding.

Example of the style (for a different paper): "Wasn'
     t there a paper looking for a threshold below
     which dietary cholesterol wouldn't harm
     arteries in primates, and they never found one?
      I think the good cholesterol dropped even
     though the total looked normal."

Requirements:
- 15-40 words, conversational English. Hedge it: "
     wasn't there a paper", "I think", "something
     about".
- Blur every specific: approximate quantities ("
     around 150", "a few hundred"), plain-language
     paraphrases of technical terms.
- Do NOT use any name the paper introduces or relies
      on: no method, model, dataset, tool, or
     benchmark names, no acronyms, no author names,
     no venue.
- Reuse at most two content words from the title;
     paraphrase the rest.
- No exact values with decimals, no identifiers, no
     URLs, no quotation marks, no figure/table
     references.
- The description must still be recognizable: keep
     the concrete setup (what was studied, on what,
     and roughly what was found).
- Always respond in English.
- If the abstract and body are too abstract to
     describe this way, respond with exactly
     NO_DISTINCT_QUERY.

Title: {{ title }}
Abstract: {{ abstract }}
Body excerpt:
{{ body }}

Respond with ONLY the description (or
     NO_DISTINCT_QUERY), on a single line. No
     explanation, no JSON.
```

Judging. `query_profile.jinja`: News and rare query profile, system prompt.

```
Classify a web search query for the Needs Met rating
      framework. Your classification is passed
     verbatim to every per-document judgement for
```

```
    this query, so all documents are judged against
     the same query interpretation. The sections
    below are shared with the per-document judge.
    References like 2/SM or 3+ are ratings on the
    judge's 0-4 Needs Met scale (0-1 FailsM, 2 SM,
    3 HM, 4 FullyM); they describe how the judge
    uses each classification. Be decisive.

## Query analysis

### MAIN vs AUXILIARY aspects classification

For each key aspect implied by the query (its
    objective, sub-questions, and the kind of
    source it calls for) we have 2 gradations:
- MAIN: Directly answers the query's primary
    objective. Foundational facts the user needs.
- AUX: Supporting details, statistics, context that
    enhances but doesn't define the answer.

Aspects describe what a document must cover, never
    the answer itself. Write "the identity of the
    author", not the author's name. Do not answer
    the query from your own knowledge: the judge
    scores documents against their own content, and
     a wrong precomputed answer would corrupt every
     judgement for this query.

List 1-3 MAIN aspects, each the single minimal fact
    a document must contain to answer the query. If
     one aspect would fully satisfy the user, list
    only that one. Dates, venues, candidates,
    outcomes, and secondary qualifiers are AUX
    unless the query explicitly asks for them.

### Query type classification

Query types:
{% include "query_types.jinja" %}
### Search operator extraction

Parse the query for explicit search operators. These
     are hard constraints the user typed
    deliberately. They are not optional concepts.

- Exact-phrase: any span inside double quotes "...".
     Each quoted span is a separate required string.

- site:domain = the document URL host must equal or
    be a subdomain of `domain`.

### Content archetype classification

Classify the query's expected content to determine
    staleness evaluation:
{% include "content_archetypes.jinja" %}
### Domain classification

Classify the query into one of these domains, or "
    other" if none fits:
- "health": Medical and health queries.
- "legal": Law, regulations, caselaw.
- "financial": Companies, tickers, markets, personal
     finance.
- "academic": Science, research, formulas.
- "tech": Software, hardware, troubleshooting.
- "news": Current events and reporting.
- "commerce": Products and shopping.
- "jobs": Job listings and hiring.

## Output

Respond with ONLY a YAML block. No text before or
    after.

'''yaml
query_type: <broad|specific|list>
query_domain: <health|legal|financial|academic|tech|
    news|commerce|jobs|other>
query_search_operators_exist: <yes|no>
query_main_aspects:
  - "<MAIN aspect>"
query_aux_aspects:
  - "<AUX aspect>"
query_content_archetype: <ephemeral|persistent|
    evergreen>
'''
```

`query_profile_input.jinja`: Query profile, user message.

```
Today's date: {{ today }}

Query: {{ query }}
```

`judgement.jinja`: News and rare relevance judge, system prompt.

```
Evaluate how well a search result document matches a
     user's query intent using the Needs Met rating
     framework (0-4 scale).

You are provided with:
- Today's date
- Query: The user's search query
- Query Analysis: precomputed classification (query
    type, domain, search-operators flag, MAIN and
    AUX aspects, content archetype)
- Document: Title, URL, optional published date,
    snippet

Adopt the Query Analysis block as ground truth and
    do not re-classify the query. Step 1 below
    defines its fields. If the block is absent,
    infer these fields from the query text yourself.
     The Query Analysis defines what to look for;
    the document itself is the only source of
    factual answers -- judge facts against the page
    's own text and the source's authority, never
    against your own knowledge of the answer. You
    may still use general knowledge to recognize
    that two names or descriptions refer to the
    same entity.

## Rating Scale

- 4 (FullyM): All MAIN aspects covered with
    substantial depth. Authoritative, comprehensive
    , directly answers the query's primary
    objective. One-stop resource.
- 3 (HM): Major MAIN aspects covered with good depth
    . Minor gaps in auxiliary details or slight
    limitations, but the reader's core need is
    still substantially satisfied.
- 2 (SM): Partial MAIN coverage, or only passing/
    superficial mention without actionable depth.
    May touch the general topic while missing the
    precise query elements. Trips at least one
    document condition.
- 1 (FailsM): Marginal relevance only. Topic-
    adjacent but fundamentally misses the query's
    relevance or provides negligible useful
    information.
- 0 (FailsM): Off-topic, broken, inaccessible,
    duplicate URL, or wholly fails the query.

Rating and label must always be consistent:
- Rating 0 = FailsM
- Rating 1 = FailsM
- Rating 2 = SM
- Rating 3 = HM
- Rating 4 = FullyM

If there is any mismatch, correct it before
    outputting.

## Step 1: Query analysis

### MAIN vs AUXILIARY aspects

For each key aspect implied by the query (its
    objective, sub-questions, and the kind of
    source it calls for) we have 2 gradations:
- MAIN: Directly answers the query's primary
    objective. Foundational facts the user needs.
- AUX: Supporting details, statistics, context that
    enhances but doesn't define the answer.

Missing AUX alone never drops a document below 3/HM
    if all MAIN aspects are adequately covered.
Additionally, queries might have hard caps that
    limit the maximum possible label.

### Query types
```

```
The query type sets how strict the content
     requirements are.

Query types:
{% include "query_types.jinja" %}
### Search operators

Parse the query for explicit search operators. These
      are hard constraints the user typed
     deliberately. They are not optional concepts.

- Exact-phrase: any span inside double quotes "..."
     must appear VERBATIM in the snippet (the
     content between `<<<SNIPPET>>>` and `<<</
     SNIPPET>>>`) or the document title (case-
     insensitive, whitespace-normalized; ignore
     surrounding punctuation). A match in the URL,
     breadcrumb, navigation links, or any other
     metadata does not count and never satisfies the
      operator. Do not assume or infer presence:
     locate the exact span in the snippet or title,
     or treat it as missing. Each quoted span is a
     separate required string.
- site:domain = the document URL host must equal or
     be a subdomain of `domain`.

When a query carries no operators, skip the operator
      gates and score normally.

### Content archetypes

The content archetype determines staleness
     evaluation:
{% include "content_archetypes.jinja" %}
## Step 2: Document analysis

### Document integrity check

Before any relevance analysis, confirm the page is
     reachable and functional:
- "fail": Fails integrity (rate 0). Error pages
     (404, 500), blank/metadata-only pages, hard
     paywalls blocking all content, pages with only
     boilerplate/navigation and no substantive
     content.
- "limit": Passes with limitation. Social media
     previews, academic abstracts behind paywalls --
      evaluate visible content, note limitation in
     reasoning. Cap at 2/SM when the whole page is
     required but unavailable.
- "pass": Passes completely. Content is informative.

### Document conditions

Three conditions constrain the maximum rating. All
     applicable conditions must hold for the
     document to qualify for 3/HM or higher. A
     condition can either pass, fail, or be not
     applicable.

Relevance condition (all query types):
- The page has to match the right topic, entity,
     jurisdiction, and domain context.
- Wrong entity = 0/FailsM no matter how good the
     content is (e.g., a Domino's page for a "Sal's
     Pizzeria Brooklyn" query).
- An under-specified reference to the same entity is
      NOT a wrong entity: a document that omits a
     middle name, an alias, or leaves a counterparty
      unnamed still passes the relevance condition
     when nothing contradicts the query's entity.
- Wrong jurisdiction = 0/FailsM for legal/financial
     queries.
- General topic match without specific query
     elements = fails the relevance condition for "
     specific" queries.

Coverage condition (all query types):
- MAIN aspects have to appear with real depth, not
     bare keywords or passing mentions.
- "Coming soon" without a date != adequate for a "
     when does X release" query.
- A title mention without explanation != adequate
     coverage.
- "broad" queries: At least one MAIN aspect with
     meaningful depth.
- "specific" queries: The specific requested answer
     must be present.
- "list" queries: ≥50% of requested items/aspects
     must be covered.
- A page that credibly overturns the query's premise
      still answers it: for "did X die", a report
     showing that X is alive or that the event
     belongs to a related entity satisfies coverage.
      Do not demand confirmation of a premise the
     document credibly refutes.
- A secondary query detail the document leaves
     unconfirmed (a date, a middle name, an unnamed
     counterparty) is a minor gap to note in
     reasoning, not a coverage failure, when the
     queried fact itself is reported.
- If condition fails, document is capped at 2/SM.

Completeness condition (for "list" queries):
- Covering <50% of requested items = capped at 2/SM.
- Covering ≥50% but <100% = eligible for 3/HM.
- Covering ~100% with depth = eligible for 4/FullyM.

Gate logic summary:
- All conditions pass -> eligible for 3/HM or 4/
     FullyM based on depth and authority
- Any condition fails -> maximum 2/SM (or 0/FailsM
     for relevance condition failures)
- See the section below for additional constraints
     for specific, dated current-event queries.

### Current event freshness gate

When the query refers to a specific, dated current
     event (e.g., "Medvedev vs Berrettini Monte-
     Carlo Masters Round of 32, April 8, 2026"),
     these constraints override the standard 2/SM
     cap. Apply strictly: topical overlap is not
     sufficient.

- Entity without event: 1/FailsM maximum. The
     document discusses the correct entity but does
     NOT address the specific event the query names.
      No exceptions. This includes: Wikipedia
     biographies, player/team profile pages, career
     statistics pages, game logs missing the
     requested date, generic relationship or
     biography articles, stats dashboards without
     the specific game, and any page that covers the
      entity generally without referencing the event
      itself.
- Wrong event: 1/FailsM maximum. The document covers
      a DIFFERENT event involving the same entity.
     Use the query to distinguish: a 1st-leg match
     article is wrong-event for a 2nd-leg query; a
     March committee markup article is wrong-event
     for an April 8 debate query; a 2024 Tommy John
     article is wrong-event for a 2026 shoulder
     injury query.
- Pre-event for result queries: 1/FailsM maximum.
     When the query seeks a result, outcome, or post
     -event status, any preview, prediction, or pre-
     event article = 1/FailsM. This includes odds/
     previews for completed games, pre-match
     analyses for result queries, and pre-
     announcement speculation for outcome queries.
- How to verify: Before rating 2+ for a specific
     dated-event query, confirm the document
     contains at least one CONCRETE detail matching
     the queried event (the specific date, score,
     quote, ruling, or named development). General
     coverage of the entity or topic without a
     concrete detail from the event = 1/FailsM.
- Scope: This gate only fires for specific, dated
     current-event queries. Otherwise, the standard
     max 2/SM for coverage condition failures
     remains unchanged.

### Staleness and temporal fit

Based on the expected content archetype, apply
     staleness checks:

"ephemeral" content:
- Strict temporal alignment is required. A page
     about 2024 performers does not serve a "2025
     performers" query.
- Breaking news: 0-48 hours = recency paramount. 2-7
      days = balance with comprehensiveness.
```

```
- Future-tense queries accept predictions/previews.
    Past-tense queries require actual results: a
    write-up published ahead of the event cannot
    say "who won".
- Recurring annual events: Previous year's info is
    acceptable until next year's is announced.

"persistent" content:
- Obsolete technology/platforms (Windows Phone,
    defunct services) = stale.
- Outdated business info (old CEO, wrong pricing,
    old contact details) = stale.
- Stale "Just launched!" banners for features that
    shipped as standard long ago = stale.

"evergreen" content:
- A recent publication date by itself doesn't make
    it stale. Only superseded factual developments,
     withdrawn methodologies, or incomplete
    timelines qualify.
- A 2016 write-up covering a settled scientific
    principle beats a shallow 2025 one.

Staleness penalties:
- Critical mismatch (wrong year event, wrong entity
    status) = 0/FailsM
- Severe (multiple outdated facts, functionally
    broken purpose) = downgrade 2 levels
- Moderate (some outdated details, minor temporal
    issues) = downgrade 1 level
- Minor/peripheral (old copyright year, minor detail
    ) = note in reasoning only

### Additional rules

Apply these principles to avoid over-penalizing high
    -quality documents:
- When a high-trust source (official government site
    , ESPN, an established news outlet) reports
    something different from what the query seems
    to expect, believe the page. Dock it only if it
     is clearly stale or a weak-authority source
    that contradicts several authorities.
- The query's criteria capture ideas, not literal
    keywords. Judge whether the page satisfies the
    underlying information need, not whether
    particular keywords appear. "battery-electric
    car" = "EV". "headcount reduction" = "layoffs".
     Give equivalent ideas full marks. The only
    exception are the spans the user placed in
    double quotes or other explicit operators,
    which must be matched literally and are not
    subject to conceptual substitution.
- For live or dynamic data (stock quotes, weather,
    sports scores), small gaps from expected values
     are fine when the source is authoritative.
    Dock only large deviations that signal
    staleness.
- Examples implied by the query are illustrative,
    not mandatory. A document covering equivalent
    examples satisfies the criterion.

### Domain-specific hard caps

Apply when the query falls into these domains:

- "health": Authority is paramount (CDC/WHO > blogs)
    . One dangerous error (unproven therapy) = 0/
    FailsM. Prioritize peer-reviewed sources and
    official guidelines.
- "legal": Jurisdiction must match. Mismatch = 0/
    FailsM. A critical mistake (wrong fees, expired
     deadlines) caps at 2/SM.
- "financial": Verify the entity (right company/
    ticker/exchange) - required. Wrong entity = 0/
    FailsM. Require concrete quantitative data, not
     vague claims.
- "academic": Core formula/principle errors = 1-2.
    Peripheral citation errors don't drop below 3/
    HM if core is correct.
- "tech": Safety notes required for risky procedures
     (editing the registry, taking hardware apart).
     Crowd-verified fixes can beat terse official
    pages.
- "news": Critical factual errors = 1-2. Separate
    reporting from speculation. Prefer established
    outlets.
- "commerce": Require an exact product match (right
    model/version/color/size). Wrong variant with
    same brand = 2/SM max. Different product line =
     0/FailsM.
- "jobs": If query specifies work modality (remote/
    telemedicine), listings must match. On-site
    listings for remote queries = significant
    downgrade.

### Search operator compliance gate (applies only
    when query has search operators)

Operators are explicit user constraints, equal in
    force to entity/jurisdiction:
- Required exact "phrase" absent from both the
    snippet and the title = 0/FailsM. A match found
     only in the URL, breadcrumb, navigation, or
    metadata does NOT satisfy the operator.
- site: host mismatch = 0/FailsM

A document may be perfectly on-topic and still fail
    this gate.

## Step 3: Output

Respond with only a YAML block. No text before or
    after.

'''yaml
reasoning: "<2-3 sentences. State the query type (
    broad/specific/list), which document conditions
     passed/failed, key evidence from the document,
     and any penalties applied (staleness, domain
    gate). Be specific about WHAT content was
    present or missing.>"
document_integrity_check: <fail|limit|pass>
document_relevance_condition: <fail|pass>
document_coverage_condition: <fail|pass>
document_completeness_condition: <fail|pass|na>
freshness_gate: <fail|pass|na>
staleness_gate: <fail|pass>
domain_gate: <fail|pass|na>
search_operators_gate: <fail|pass|na>
rating: <0-4>
label: <FailsM|SM|HM|FullyM>
'''

Gate field rules:
- document_relevance_condition and
    document_coverage_condition apply to every
    query: always pass or fail.
- document_completeness_condition: na unless the
    query type is "list".
- freshness_gate: na unless the query targets a
    specific, dated current event.
- staleness_gate: fail only when a staleness penalty
     applies; pass otherwise.
- domain_gate: na when the query domain is "other".
- search_operators_gate: na when the query has no
    search operators.

Output rules:
- YAML only. No markdown fences, no commentary, no
    extra fields.
- `reasoning` must be 2-3 sentences grounded in
    document evidence.
- Rating and label must be consistent (see rule
    above).
```

`judgement_input.jinja`: Judge, user message.

```
Today's date: {{ today }}

Query: {{ query }}

{% if profile -%}
Query Analysis (precomputed -- adopt as ground truth
    , do not re-classify):
- Query type: {{ profile.query_type }}
- Domain: {{ profile.query_domain }}
- Search operators exist: {{ "yes" if profile.
    query_search_operators_exist else "no" }}
- Content archetype: {{ profile.
    query_content_archetype }}
{% if profile.query_main_aspects -%}
- MAIN aspects:
```

```
{% for aspect in profile.query_main_aspects -%}
{{ "  - " ~ aspect }}
{% endfor -%}
{% endif -%}
{% if profile.query_aux_aspects -%}
- AUX aspects:
{% for aspect in profile.query_aux_aspects -%}
{{ "  - " ~ aspect }}
{% endfor -%}
{% endif %}
{% endif -%}
Document to evaluate:
- Title: {{ title }}
- URL: {{ url }}
{% if published -%}
- Published: {{ published }}
{% endif -%}
{% if content -%}
- Page Snippet:
<<<SNIPPET>>>
{{ content }}
<<</SNIPPET>>>
{% endif -%}
```

query_types.jinja: Included by the profile and judge prompts.

```
- "broad": General or exploratory topic queries ("
     what is inflation", "zendaya", "origins of jazz
     ").
  - MAIN = foundational defining facts. AUX =
     statistics, supplementary details.
  - The page has to cover at least one main aspect
     in real depth to reach 3+.
- "specific": Queries after a single closed answer
     ("when X launches", "who won the final", "how
     much Y is worth").
  - MAIN = the specific requested fact. Nothing else
     substitutes.
  - Hard cap: no specific answer present = capped at
     2/SM, regardless of how much related context
     is provided. A page that credibly overturns the
     query's false premise counts as providing the
     answer.
- "list": Comparisons, multi-item requests, pros-and
     -cons ("A vs B", "best 10 albums", "upsides and
     downsides of Z").
  - MAIN = coverage of all requested items/aspects.
  - Hard cap: Covering <50% of requested items =
     capped at 2/SM.
```

content_archetypes.jinja: Included by the profile and judge prompts.

```
- "ephemeral": Sports scores, weather, breaking news
     , stock prices, event schedules. Value decays
     rapidly.
- "persistent": Official websites, product pages,
     how-to guides for software. URL always relevant
     but content freshness matters.
- "evergreen": Historical facts, scientific
     principles, definitions, recipes. Largely time-
     insensitive.
```

answer_match.jinja: Finance LLM backstop for non-money fields.

```
You are verifying search results for a company fact-
     lookup benchmark.

Search query: {{ query }}
Fact being looked up: {{ field }}
Gold answer: {{ value }}
{% if aliases -%}
Also acceptable forms: {{ aliases | join("; ") }}
{% endif %}
Search result:
- Title: {{ title }}
- URL: {{ url }}
- Snippet: {{ snippet }}

Does this search result state the gold answer? Any
     equivalent form counts --
different wording, formatting, units, reasonable
     rounding (e.g. "$416.2B" for
416161000000), nicknames, or the answer appearing in
     the result URL. The result
must actually contain the fact; being about the
     right company or topic is not
enough. If the result states a clearly different
     value, answer no.

Answer with a single word: yes or no.
```